

In HIS name

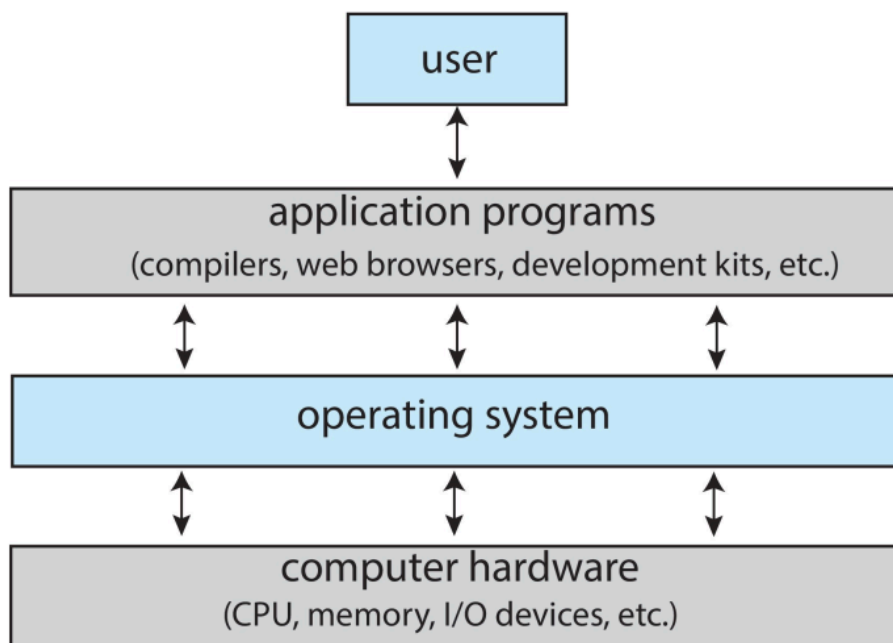
OS

Chapter 1

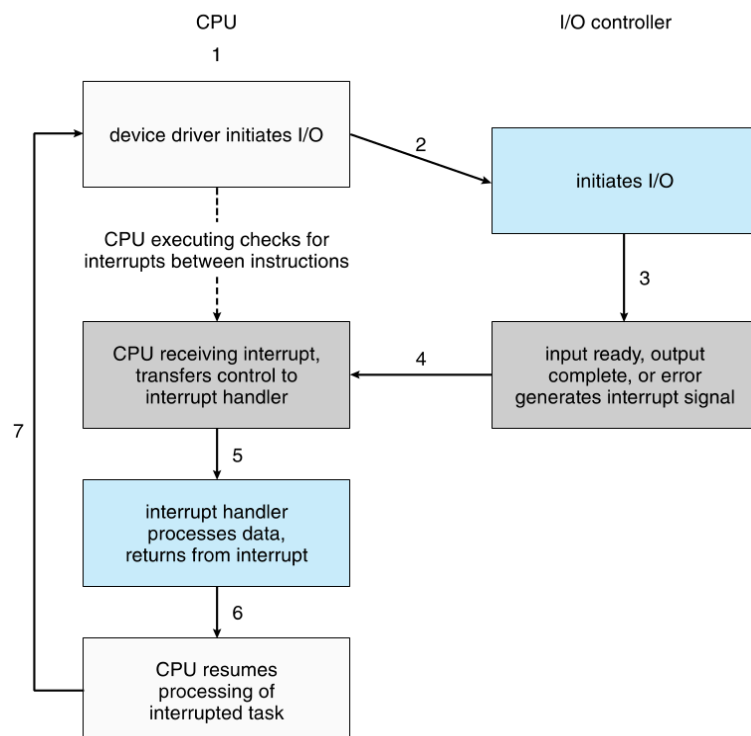
→ Operating system:

- ◆ CPU, memory, I/O devices
- ◆ Controls and coordinates use of hardware among various
- ◆ One purpose of OS is to hide peculiarities of hardware devices from the user

→ Operating system is a resource allocator and control program making efficient use of HW and managing execution of user programs



- “The one program running at all times on the computer” is the kernel, part of the operating system
- Others are:
 - ◆ system program (ships with the operating system, but not part of the kernel) , or
 - ◆ an application program, all programs not associated with the operating system
- Shares memory : One or more CPUs, device controllers connect through common bus providing access to shared memory
- Interrupt:
 - ◆ Device controller informs CPU that it has finished its operation by causing an interrupt
 - ◆ Interrupt transfers control to the interrupt service routine generally, through the interrupt vector, which contains the addresses of all the service routines
 - ◆ Interrupt architecture must save the address of the interrupted instruction
- A **trap** or **exception** is a software-generated interrupt caused either by an **error** or a **user request**
- An operating system is **interrupt driven**



→ Memory:

- ◆ Random access
- ◆ Typically random-access memory in the form of Dynamic Random-access Memory (DRAM)
- ◆ Typically volatile

→ Secondary storage : extension of main memory that provides **large nonvolatile storage capacity**

→ Hard Disk Drives (HDD) : rigid metal covered with magnetic recording material

→ Disk surface is logically divided into **tracks**, which are subdivided into **sectors**

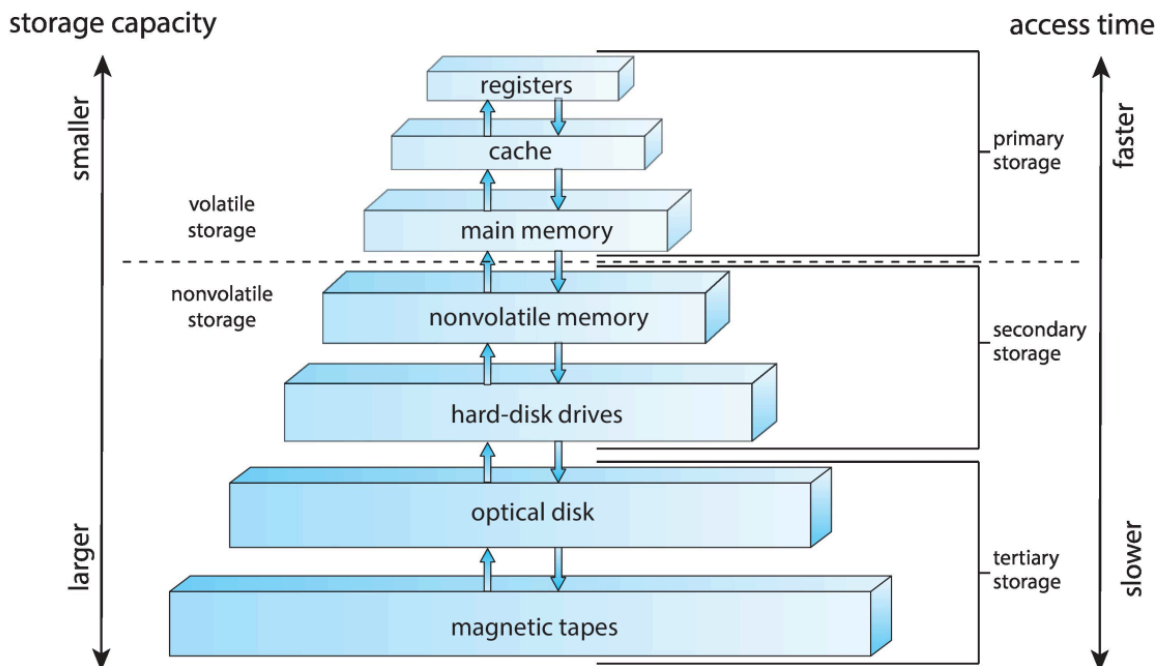
→ The disk **controller** determines the **logical interaction between the device and the computer**

→ Caching : copying information into faster storage system; main memory can be viewed as a cache for secondary storage

→ Device Driver : for each device controller to manage I/O Provides uniform **interface between controller and kernel**

→ Storage systems organized in hierarchy:

- ◆ Speed Volatility Cost






→ Direct memory access:

- ◆ Used for high-speed I/O devices able to transmit information at close to memory speeds
- ◆ Device controller transfers blocks of data from buffer storage directly to main memory without CPU intervention
- ◆ Only one interrupt is generated per block, rather than the one interrupt per byte

→ Multiprocess

- ◆ Increased throughput
- ◆ Economy of scale
- ◆ Increased reliability – fault tolerance
- ◆ Types:
 - **Asymmetric** Multiprocessing – each processor is assigned a specific task.
 - **Symmetric** Multiprocessing – each processor performs all tasks






- Like multiprocessor systems, but multiple systems working together Usually sharing storage via a **storage-area network (SAN)**
- Provides a high-availability service which survives failures
 - ◆ **Asymmetric** clustering has one machine in hot-standby mode
 - ◆ **Symmetric** clustering has multiple nodes running applications, monitoring each other
- Some clusters are for high-performance computing (HPC)
 - ◆ Applications must be written to use parallelization
- Some have distributed lock manager **(DLM) to avoid conflicting operations**
- Bootstrap program : simple code to initialize the system, load the kernel
 - ◆ Kernel loads

- ◆ Starts system daemons (services provided outside of the kernel)
- Hardware interrupt by one of the **devices**
- Software interrupt (exception or trap):
 - ◆ Software error (e.g., division by zero)
 - ◆ Request for operating system service – system call
 - ◆ Other process problems include infinite loop, processes modifying each other or the operating system

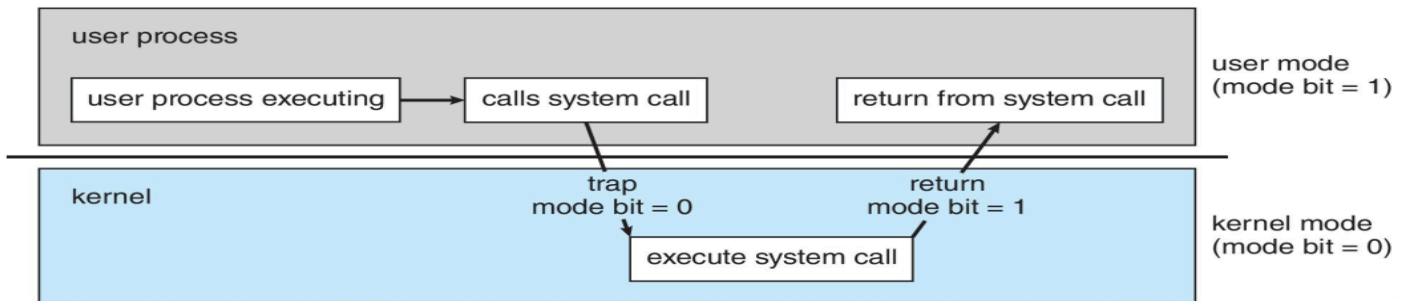
- **Multiprogramming (Batch system)** : needed for efficiency
 - ◆ Single user cannot keep CPU and I/O devices busy at all times
 - ◆ A subset of total jobs in system is kept in memory
 - ◆ Multiprogramming organizes jobs (code and data) so CPU always has one to execute
 - ◆ One job selected and run via job scheduling
 - ◆ When it has to wait (for I/O for example), OS switches to another job



- Timesharing (multitasking) is logical extension in which CPU switches jobs **so frequently** that users can **interact with each job while it is running**,
 - ◆ Response time should be **< 1 second**
 - ◆ if several jobs ready to run at the same time -> CPU scheduling
 - ◆ Virtual memory allows execution of processes not completely in memory
- Dual-mode operation:
 - ◆ allows OS to protect itself and other system components
 - ◆ Provides ability to distinguish when system is running user code or kernel code
 - ◆ User mode and kernel mode
 - ◆ Mode bit provided by hardware
 - ◆ Some instructions designated as privileged, only executable in kernel mode

→ multi-mode operations:

- ◆ i.e. virtual machine manager (VMM) mode for guest VMs



→ Timer to prevent infinite loop:

- ◆ Keep a counter that is decremented by the physical clock
- ◆ Operating system set the counter (privileged instruction)
- ◆ When counter zero generate an interrupt
- ◆ Set up before scheduling process to regain control or terminate program that exceeds allotted time

→ Multi-threaded process has one program counter per thread

→ The operating system is responsible for the following activities in connection with process management:

- ◆ **Creating and deleting** both user and system processes
- ◆ Providing mechanisms for **process synchronization**
- ◆ **Suspending and resuming** processes
- ◆ Providing mechanisms for **process communication**
- ◆ Providing mechanisms for **deadlock handling**

→ Faster storage (cache) checked first to determine if

→ information is there

- ◆ If it is, information used directly from the cache (fast)
- ◆ If not, data copied to cache and used there

→ Register -> compiler

→ Cache -> hardware

→ Memory, disk, tertiary storage -> OS

→ Multiprocessor environment must provide **cache coherency** in

hardware such that all CPUs have the most recent value in their Cache

→ I/O subsystem responsible for

◆ Memory management of I/O including

- buffering (storing data temporarily while it is being transferred),
- caching (storing parts of data in faster storage for performance),
- spooling (the overlapping of output of one job with input of other jobs)

◆ General device-driver interface

◆ Drivers for specific hardware devices

→ Protection : any mechanism for controlling access of processes or users to resources defined by the OS

→ Security : defense of the system against internal and external attacks

◆ Huge range, including denial-of-service, worms, viruses, identity theft, theft of service

→ Virtualization

◆ Allows operating systems to run applications within other OSes

Emulation : used when source CPU type different from target

→ type (i.e. PowerPC to Intel x86)

◆ Generally slowest method

◆ When computer language not compiled to native code
– Interpretation

→ Virtualization : OS natively compiled for CPU, running guest OSes also natively compiled

◆ Consider VMware running WinXP guests, each running applications, all on native WinXP host OS

◆ VMM (virtual machine Manager) provides virtualization services

→ Benefits of Virtualizations:

◆ Developing apps for multiple OSes without having multiple systems


- ◆ QA testing applications without having multiple systems
Executing and managing compute environments within data centers



- Distributed computing
 - ◆ Collection of separate, possibly heterogeneous, **systems networked together**
- Compute-server system “
 - ◆ provides an interface to client to **request services** (i.e., database)
- File-server system :
 - ◆ provides interface for clients to **store and retrieve files**

→ Another model of distributed system:

- ◆ P2P does not distinguish clients and servers
- ◆ Instead all nodes are considered peers
- ◆ May each act as client, server or both
- ◆ Broadcast request for service and respond to requests for service via discovery protocol
- ◆ Examples include Napster and Gnutella, Voice over IP (VoIP) such as Skype



→ Delivers computing, storage, even apps as a service across a network


→ Logical extension of virtualization because it uses virtualization as the base for its functionality.

- ◆ Public cloud : available via Internet to anyone willing to pay
- ◆ Hybrid cloud – includes both public and private cloud components
- ◆ Private cloud – run by a company for the company's own use

→ Software as a Service (SaaS) – one or more applications available via the Internet (i.e., word processor)

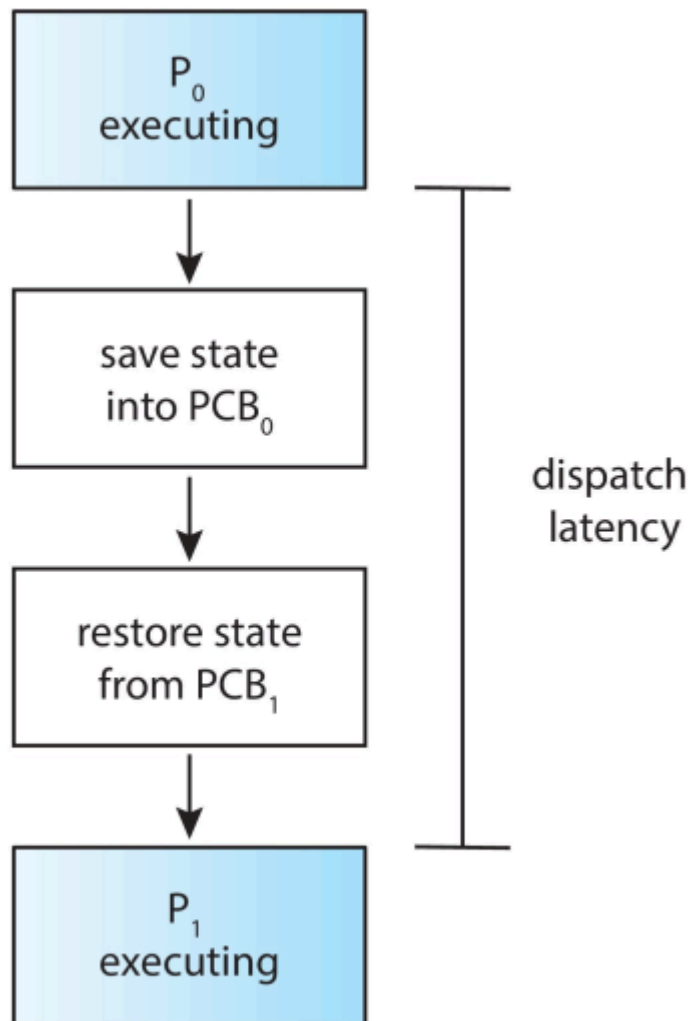
→ Platform as a Service (PaaS) – software stack ready for application use via the Internet (i.e., a database server)

- Infrastructure as a Service (IaaS) – servers or storage available over Internet (i.e., storage available for backup use)
- Load balancers spread traffic across multiple applications



Q5

- Dispatcher module:
 - ◆ switching context
 - ◆ switching to user mode
- Dispatch latency : time it takes for the dispatcher to stop one process and start another running

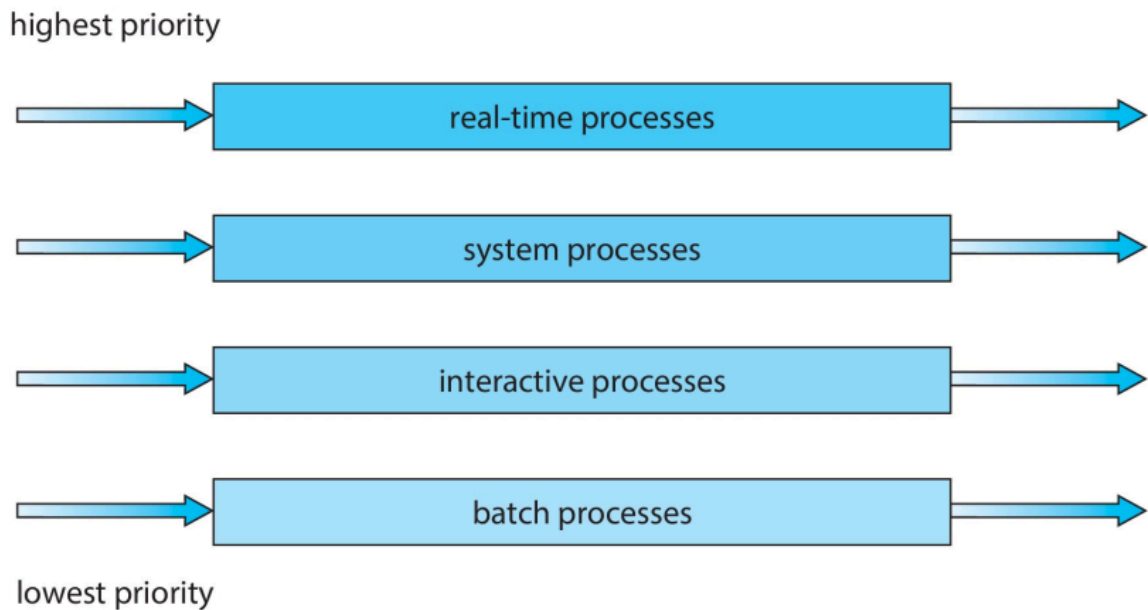


- CPU utilization (Max)– keep the CPU as busy as possible
- Throughput (Max)– # of processes that complete their execution per time unit
- Turnaround time (Min)– amount of time to execute a particular process
- Waiting time (Min)– amount of time a process has been waiting in the ready queue
- Response time (Min)– amount of time it takes from when request was submitted until the first response is produced, not output (for time-sharing environment)

- FCFS = FIFO
- **Convoy effect** - short process behind long process
- Consider one CPU-bound and many I/O-bound processes

- SJF (Shortest job first) is optimal – gives minimum average waiting time
- The difficulty is knowing the length of the next CPU request
- t_n = actual length of n th CPU burst
- 2. τ_{n+1} = predicted value for the next CPU burst
- 3. α , $0 \leq \alpha \leq 1$
- 4. Define : $\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$.
- Commonly, α set to $\frac{1}{2}$
- Preemptive version called **shortest-remaining-time-first**
- Shortest-remaining-time-first = preemptive sjf (on arrival of processes)
- Round Robin
 - ◆ q large \Rightarrow FIFO
 - ◆ q small $\Rightarrow q$ must be large with respect to context switch, otherwise overhead is too high
 - ◆ Typically, higher average waiting time
 - ◆ Typically, higher average turnaround than SJF
 - ◆ q should be large compared to context switch time
 - ◆ q usually 10ms to 100ms, context switch < 10 usec
 - ◆ 80% of CPU bursts should be shorter than q
- Priority Scheduling = (smallest integer \equiv highest priority)
 - ◆ Problem \equiv Starvation – low priority processes may never execute
 - ◆ Solution \equiv Aging – as time progresses increase the priority of the process
- Priority with RR = Processes with the same priority run round-robin

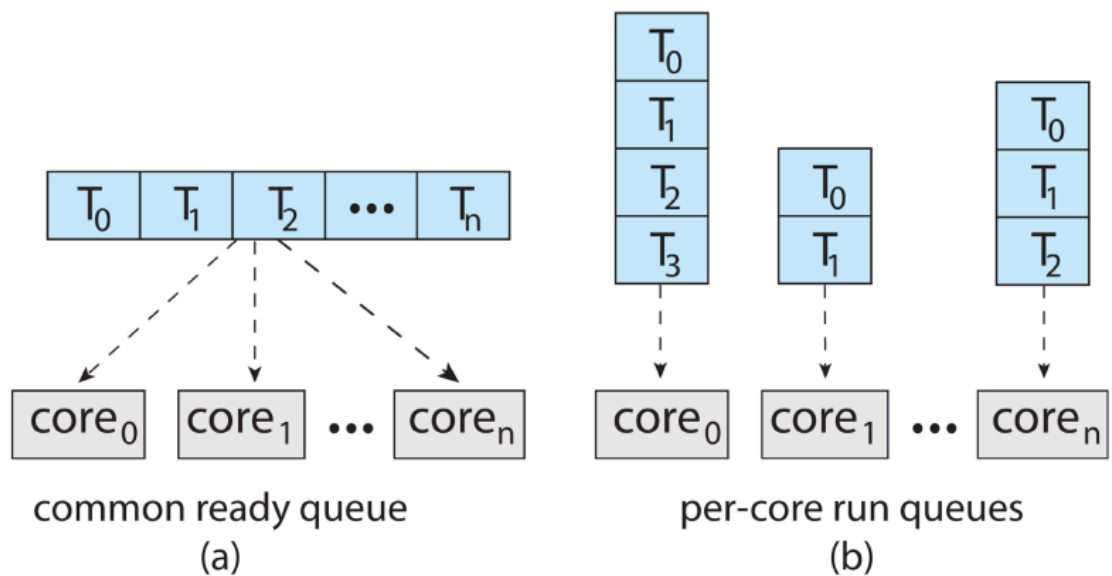
→ MLFQ = have separate queues for each priority.




- ◆ number of queues
- ◆ method used to determine when to upgrade a process
- ◆ method used to determine which queue a process will enter when that process needs service
- ◆ scheduling algorithms for each queue
- ◆ method used to determine when to demote a process

→ **Symmetric multiprocessing (SMP)** is where each processor is self scheduling.

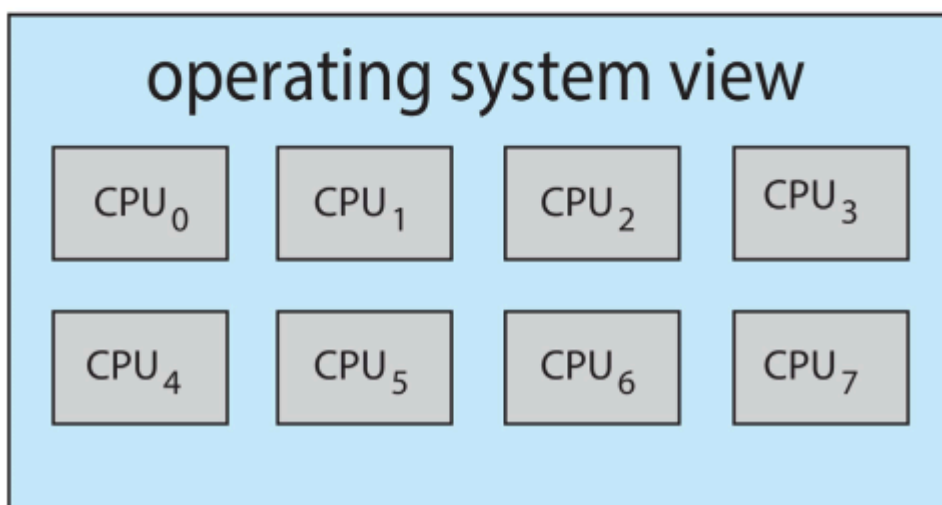
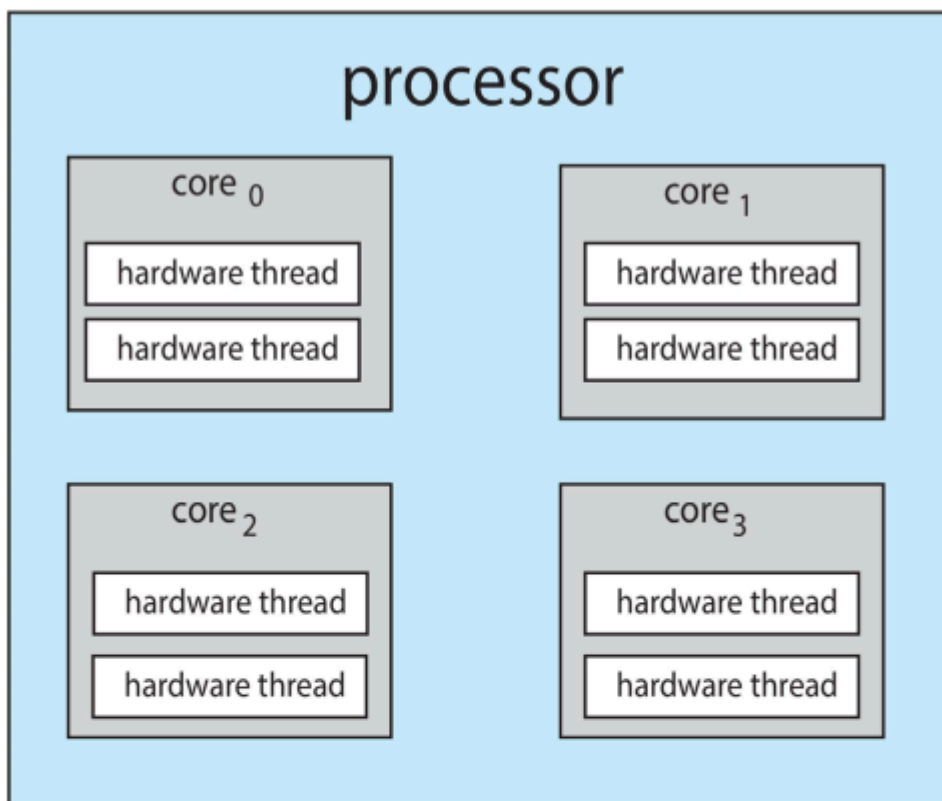
- ◆ All threads may be in a common ready queue (a)
- ◆ Each processor may have its own private queue of threads (b)



- Place multiple processor cores on same physical chip
- Faster and consumes less power
- Multiple threads per core also growing
- If one thread has a memory stall, switch to another thread!




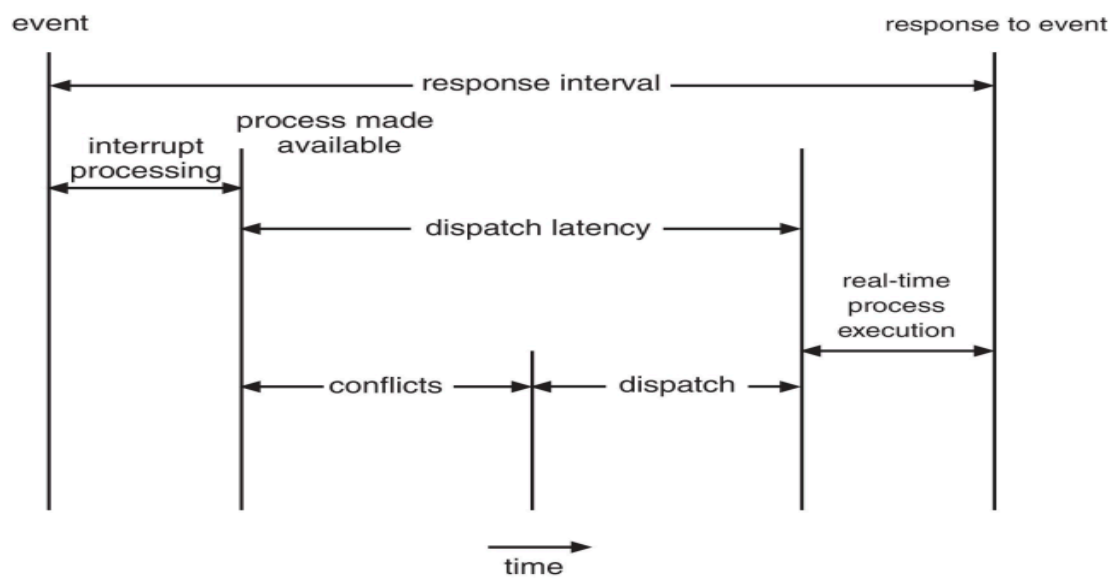
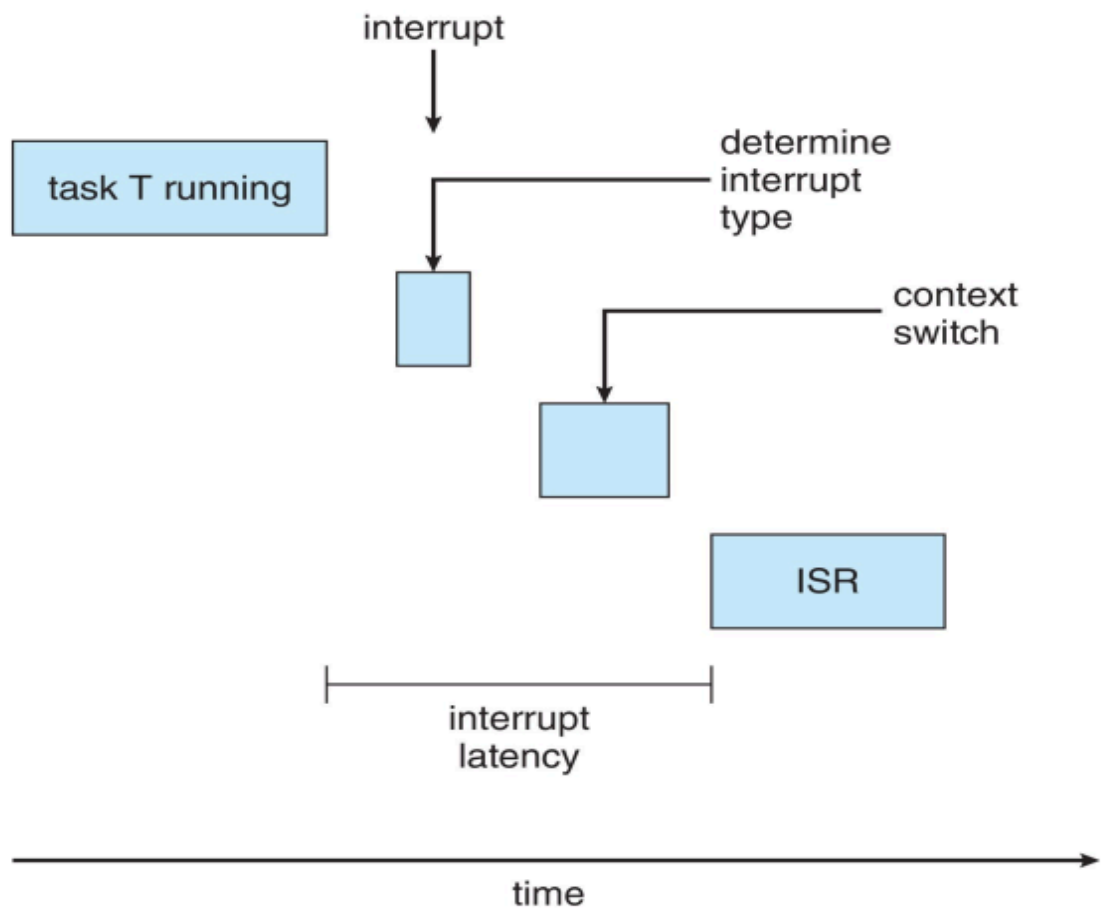
- Chip-multithreading (CMT)
- assigns each core multiple hardware threads. (hyperthreading.)
- On a quad-core system with 2 hardware threads per core, the operating system sees 8 logical processors.



- If **SMP**, need to keep all CPUs loaded for efficiency
- **Push migration** – **periodic task** checks load on each processor, and if found **pushes task from overloaded CPU to other CPUs**
- **Load balancing** attempts to **keep workload evenly distributed**
- **Pull migration** – **idle processors pulls waiting task from busy processor**

- Thread loses the contents of what it had in the cache of the processor it was moved off of.
- **Soft affinity** – the operating system attempts to keep a thread running on the same processor, but no guarantees.
- **Hard affinity** – allows a process to specify a set of processors it may run on.
- **NUMA-aware** OSES will assign memory closest to the CPU the thread is running on.
- **Real time**
 - ◆ **Soft real-time** systems – Critical real-time tasks have the highest priority, but no guarantee as to when tasks will be scheduled
 - ◆ **Hard real-time** systems – task must be serviced by its deadline
- Event latency – the amount of time that elapses from when an event occurs to when it is serviced.
- 1. **Interrupt latency** – time from arrival of interrupt to start of routine that services interrupt
- 2. **Dispatch latency** – time for schedule to take current process off CPU and switch to another
- 1. Preemption of any process running in kernel mode
- 2. Release by low- priority process of resources needed by high- priority processes





- ◆ preemptive, priority-based scheduling
- ◆ only guarantees soft real-time
- ◆ hard real-time must also provide ability to meet deadlines
- Rate-Monotonic
 - ◆ Rate of periodic task is $1/p$
 - ◆ Shorter periods = higher priority;
- EDF (Earliest deadline first)
 - ◆ the earlier the deadline, the higher the priority;

→ Algorithm Evaluation

- ◆ How to select a CPU-scheduling algorithm for an OS?
- Deterministic modeling:
 - ◆ Type of analytic evaluation
 - ◆ Takes a particular predetermined workload and defines the
 - ◆ performance of each algorithm for that workload
 - ◆ Simple and fast, but requires exact numbers for input, applies only to those inputs

→ Queueing Models:

- ◆ Describes the arrival of processes, and CPU and I/O bursts probabilistically
- ◆ Commonly exponential, and described by mean
- ◆ Computes average throughput, utilization, waiting time, utilization, average queue length, average wait time, etc.
- ◆ Knowing arrival rates and service rates

→


- n = average queue length
- λ = average arrival rate into queue
- W = average waiting time in queue
- Little's law – in the steady state, processes leaving queue must equal processes arriving, thus:

→ $n = \lambda \times W$

→ Simulations

- ◆ Queueing models limited
- ◆ Simulations more accurate
- ◆ Programmed model of computer system
- ◆ Gather statistics, indicating algorithm performance
- ◆ Clock is a variable
- ◆ Data to drive simulation gathered via
 - Random number generator according to probabilities
 - Distributions defined mathematically or empirically
 - Trace tapes record sequences of real events in real systems
- ◆ Even simulations have limited accuracy
- ◆ Just implement new scheduler and test in real systems
 - High cost, high risk
 - Environments vary

- ◆ Most flexible schedulers can be modified per-site or



per-system